



December 6, 2025

## Learning Goal

Understand the vanishing gradient problem caused by sigmoid saturation.

---

This slide establishes the learning objective for this topic

## Key Concept (1/2)

The sigmoid function **saturates** at extreme values. When  $z$  is very large (e.g.,  $z \geq 3$ ) or very small ( $z \leq -3$ ), the sigmoid output is nearly 1 or 0, and barely changes regardless of input changes.

This creates the **vanishing gradient problem**. During backpropagation, gradients flow backward through the network. In saturated regions, gradients become extremely small (near zero), so weights barely update. The network stops learning.

---

Understanding this concept is crucial for neural network fundamentals

## Key Concept (2/2)




The **maximum gradient** of sigmoid is only 0.25 (at  $z = 0$ ). In deep networks, multiplying many small gradients results in effectively zero gradient for early layers - they learn very slowly or not at all.

**ReLU** solves this by having a constant gradient of 1 for positive inputs. However, ReLU has its own issue: neurons can "die" if  $z$  is always negative (gradient = 0).

---

Understanding this concept is crucial for neural network fundamentals

# Visualization



Visual representations help solidify abstract concepts

## Key Formula

**Sigmoid and its derivative:**

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

$$\sigma'(z) = \sigma(z)(1 - \sigma(z))$$

**Maximum derivative:**

$$\max(\sigma'(z)) = \sigma(0)(1 - \sigma(0)) = 0.5 \times 0.5 = 0.25$$

**ReLU and its derivative:**

$$\text{ReLU}(z) = \max(0, z)$$

$$\text{ReLU}'(z) = \begin{cases} 1 & z > 0 \\ 0 & z \leq 0 \end{cases}$$

---

Mathematical formalization provides precision

Imagine pushing a ball on different surfaces:

- **Sigmoid at extremes:** Like pushing a ball on a flat plateau. No matter how hard you push, it barely moves. (Small gradient = little learning)
- **Sigmoid at center:** Gentle slope, ball rolls smoothly. (Gradient = 0.25)
- **ReLU positive region:** Constant slope, ball rolls consistently. (Gradient = 1)
- **ReLU negative region:** Another flat plateau, ball doesn't move.

For efficient learning, we want consistent, non-vanishing gradients throughout.

---

Intuitive explanations bridge theory and practice

## Practice Problem 1

### Problem 1

Calculate the sigmoid derivative at  $z = 5$ . How does it compare to the maximum derivative?

### Solution

First, calculate sigmoid:

$$\sigma(5) = \frac{1}{1 + e^{-5}} = \frac{1}{1 + 0.0067} = 0.9933$$

Then, calculate derivative:

$$\sigma'(5) = \sigma(5)(1 - \sigma(5)) = 0.9933 \times 0.0067 = 0.0067$$

**Comparison to maximum:**

$$\frac{0.0067}{0.25} = 0.027 = 2.7\%$$

At  $z = 5$ , the gradient is only **2.7% of maximum**. Learning is extremely slow in this saturated region.

---

Practice problems reinforce understanding

## Practice Problem 2

### Problem 2

A network has 5 layers, each with sigmoid activation. If each layer's gradient is 0.2, what is the total gradient that reaches the first layer?

### Solution

Gradients multiply through layers (chain rule):

$$\begin{aligned}\text{Total gradient} &= 0.2 \times 0.2 \times 0.2 \times 0.2 \times 0.2 = 0.2^5 \\ &= 0.00032\end{aligned}$$

The first layer receives a gradient of only 0.00032 - practically zero!

With learning rate 0.1:

$$\Delta w = 0.1 \times 0.00032 = 0.000032$$

Weight updates are negligible. The first layer cannot learn effectively.

**This is the vanishing gradient problem.**

---

Practice problems reinforce understanding

## Key Takeaways

- Sigmoid saturates at extreme values (gradient near 0)
- Maximum sigmoid gradient is only 0.25
- Vanishing gradients prevent learning in deep networks
- ReLU has constant gradient = 1 for positive inputs
- Modern networks use ReLU (or variants) for hidden layers

---

These key points summarize the essential learnings