



December 6, 2025

## Learning Goal

Visualize the error surface that neural networks navigate during training.

---

**This slide establishes the learning objective for this topic**

The **loss function** measures how wrong our predictions are. For each possible combination of weights, we can calculate the total error across all training examples. If we plot this error as a function of weights, we get the **loss landscape** - a surface with hills (high error) and valleys (low error).

---

Understanding this concept is crucial for neural network fundamentals

## Key Concept (2/3)

Training a neural network is like finding the lowest point in this landscape. The **global minimum** is the lowest point overall - the best possible weights. Local minima are lower points surrounded by higher ground, but not the absolute lowest.

With only two weights, we can visualize this as a 3D surface. With millions of weights (typical for modern networks), the landscape exists in millions of dimensions - impossible to visualize but mathematically identical.

---

Understanding this concept is crucial for neural network fundamentals



## Key Concept (3/3)

The shape of the loss landscape determines how easy or hard training will be. Smooth landscapes with few local minima are easier to optimize than rugged landscapes with many traps.

---

**Understanding this concept is crucial for neural network fundamentals**

# Visualization



Goal: Find the weights that minimize the loss  
(The red star shows the optimal solution)

Visual representations help solidify abstract concepts

**Binary Cross-Entropy Loss** (for classification):

$$L = -\frac{1}{N} \sum_{i=1}^N [y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)]$$

Where: - **N** = number of training examples - **y<sub>i</sub>** = true label (0 or 1) - **y-hat<sub>i</sub>** = predicted probability - **L** = loss value (lower is better)

---

Mathematical formalization provides precision

Imagine you're blindfolded on a hilly terrain, trying to find the lowest valley. You can only feel the slope directly beneath your feet. The loss landscape is like this terrain, where:

- Your position = current weight values - Altitude = prediction error (loss) - Goal = find the lowest altitude (minimum loss)

The challenge: you might get stuck in a small dip (local minimum) without realizing there's a deeper valley nearby.

Training strategies like momentum and learning rate schedules help escape such traps.

---

**Intuitive explanations bridge theory and practice**

## Practice Problem 1

### Problem 1

Calculate the binary cross-entropy loss for a single prediction where the true label is  $y = 1$  and the predicted probability is  $\hat{y} = 0.9$ .

### Solution

$$\begin{aligned}L &= -[y \log(\hat{y}) + (1 - y) \log(1 - \hat{y})] \\L &= -[1 \cdot \log(0.9) + (1 - 1) \cdot \log(1 - 0.9)] \\L &= -[\log(0.9) + 0] \\L &= -(-0.105) = 0.105\end{aligned}$$

The loss is **0.105**. This is low because our prediction (90%) closely matches the true label (1).

---

Practice problems reinforce understanding

## Practice Problem 2

### Problem 2

Now calculate the loss when  $y = 1$  but  $\hat{y} = 0.1$  (a bad prediction).

### Solution

$$L = -[1 \cdot \log(0.1) + 0 \cdot \log(0.9)]$$

$$L = -[\log(0.1)]$$

$$L = -(-2.303) = 2.303$$

The loss is **2.303** - much higher than before! This penalizes confident wrong predictions severely.

Comparison: - Good prediction ( $\hat{y}=0.9$ ): Loss = 0.105 - Bad prediction ( $\hat{y}=0.1$ ): Loss = 2.303 (22x worse)

---

Practice problems reinforce understanding

## Key Takeaways

- Loss functions quantify prediction error
- The loss landscape shows error as a function of weights
- Training seeks to find the minimum loss (optimal weights)
- Cross-entropy severely penalizes confident wrong predictions
- Landscape shape affects how easily we can find good solutions

---

These key points summarize the essential learnings